00:00 Hey everybody, I'm Jamie. Welcome back to Teachers Tech. So, in my last Claude Code video,
00:05 we did something pretty cool. We built a full bookmark dashboard app just by typing in plain
00:10 English. No coding experience needed. If you haven't seen that one, I'll link it right up
00:15 here and down below in the description. Definitely go watch that one first because today's video
00:20 builds on everything we covered there. But today, we're going to go way beyond building apps. Today,
00:25 we're going to be building an AI agent. Let me show you what I mean. Are you seeing this?
00:30 I typed one sentence and it started thinking on its own. It broke the task into steps. It asked
00:35 me follow-up questions. It went out to gather information, organized everything, and produced
00:40 a full report all by itself. Last time, Cloud Code did what we told it to do. This time, we're
00:46 going to teach it to figure out things on its own. That's called an agentic workflow, and honestly,
00:51 it's the most useful thing I've learned in AI this year. Oh, and we're doing the whole thing inside
00:55 VS Code this time. way more comfortable than the raw terminal like last time. Let's get into
01:00 it. All right, before we start building anything, let's talk about what an agentic workflow actually
01:08 is because it sounds way more complicated than it really is. Think back to the last video. We told
01:13 Claude Code exactly what to build. We said, "Make me a bookmark dashboard with search dark mode."
01:19 We gave it specific instructions and it followed them. That was powerful, but we're still doing
01:24 all the thinking. We were directing every step. An agentic workflow flips that around. Let me
01:29 break it down into three levels so you can see the differences. Level one is chat. This is the most
01:35 basic way people use AI. You ask a question, you get an answer. One and done. That's your standard
01:40 chat GPT or clawed conversation. Most people never get past this level. Level two is building. This

01:48 is what we did in the last video. You tell
the AI what you want to create. You guide the
01:52 process step by step and it writes the code for
you. You're the director, the AI is the builder.
01:57 It's way more powerful than just chatting, but
you're still making all the decisions. Level three
02:02 is agentic. And this is where things get really
interesting. Instead of telling the AI what to
02:07 do step by step, you describe the goal. You say,
"Here's what I need to accomplish." And the AI
02:13 figures out the steps on its own. It makes a plan,
picks its tools, adapts when things don't go as
02:18 expected, and it will even ask you questions when
it's not sure about something. That's what we're
02:23 doing today. Here's the way I think about it. It's
like the difference between micromanaging a brand
02:28 new employee versus trusting an experienced one.
In the last video, we were the manager saying,
02:33 "Do this, then this, then this." Today,
we're saying, "Here's the goal. Here's
02:38 what you have access to. Figure it out." That's a
completely different relationship with the AI. So
02:44 what specifically makes something agentic? Three
things. First, it follows a process. It's not just
02:50 answering a single question. It's working through
a whole series of steps to get a result. Second,
02:55 it makes decisions. It chooses what to do based
on what it finds along the way. If one approach
03:00 isn't working, it pivots. If it finds something
unexpected, it adjusts. Third, it asks before it
03:06 assumes. Instead of guessing what you want, a good
aenic workflow will ask you clarifying questions
03:11 first. That one change alone makes the output
10 times better. And the use cases for this are
03:17 endless. You could use it to research a topic
across multiple sources and compile a report.
03:22 Analyze a folder of files and organize them by
category. Draft personalize emails for a whole
03:27 list of contacts. Compare competitors and produce
a sidebyside breakdown. The pattern is always

03:32 the same. You give it the goal and it does the
work. All right, let's get our tool set up. Now,
03:37 last time we worked entirely in the terminal and
that works, but VS Code gives us a much better
03:43 experience for what we're going to do today. You
can see your files, your code, and cloud code all
03:48 in the same place. When the agent is creating and
editing multiple files, which it will be, having
03:54 that visual file explorer makes a huge difference.
So, let's get this set up. This only takes about
03:59 2 minutes. If you don't have Visual Studio Code
on your computer, I'll put the link to this site
04:04 down below in the description. Then you can pick
what operating system you're working on. Today I'm
04:09 going to be on Windows, so I'm going to go ahead
and choose this and download this for my machine.
04:15 Once you have it downloaded, go ahead open it up.
Where I want you to go is right here to extensions
04:22 and you're going to type in just claude code
up here. And you're going to see that we can
04:26 install cloud code here. So I'm just going to go
ahead and click on this. I'm going to trust the
04:31 publisher and this install. All right, it's all
ready to go. We have it installed. I can see the
04:36 cloud code. I'm going to click on it to open. I'm
going to go and mark this as done. I'm just going
04:42 to go and close these extensions over here. I'll
just click on it again to give myself more room.
04:47 Notice you can uh stretch these uh to whatever
you want that works best for your preferences.
04:53 If you already installed Cloud Code from the
last video, you're all set. The extension uses
04:57 the same installation. If you haven't installed it
yet, go watch that first video because we covered
05:02 the whole setup there. Let me give you a quick
orientation of VS Code. Over here on the left,
05:07 we've got our file explorer. This is where you're
going to see every file the agent creates in real
05:13 time. Folders appearing, documents being written.
It's really satisfying to watch. Over here, we

05:19 have our cloud code panel. This is where you talk
to the agent. You go down here, type your prompts,
05:24 you see what it's thinking, and you watch it work.
Now, let's create our project folder. I'm going to
05:29 go ahead and create a new folder where we're going
to be working out of. I'm just going to click open
05:33 folder. And I haven't created one yet. If you know
what folder you want, you can go ahead and choose
05:37 that one. I'm going to go in my documents here.
And I'm just going to create a brand new folder.
05:43 And I'm going to go call it my first agent.
Select folder. I'm just going to go back and
05:48 click on clawed code here and new session just to
open that back up. And if I click on the explorer,
05:55 you can see here's the folder that I just created.
Now, there's a couple of features in VS Code that
06:01 are going to be really important for you today.
First, you can see the files updating real time as
06:06 the agent works. Watch the file explorer. When the
agent creates something, it just appears. Second,
06:12 when the agent plans a multi-step task, you'll
see a to-do list pop up showing you exactly
06:17 what's going to do and where in the process.
And here's the big one, planning mode. Press
06:22 shift plus tab and you'll toggle into planning
mode. In this mode, Claude Code thinks and plans
06:28 but doesn't actually change any files. It's like
a dry run. We're going to use this a lot today
06:32 because the key to aenic workflows is planning
before building. All right, our environment is
06:37 ready. Now, let's set up one more thing that's
really going to make all the difference. Okay,
06:42 this is something I didn't cover in the first
video and it completely changes how claude code
06:46 works for you. It's a file called claude.md.
And I'm not exaggerating when I say this is
06:52 the single most important thing you can set up.
So, what is claw.md? It's a simple markdown file
06:57 that you put in your project folder. Claude code
reads it automatically every time it starts up.

07:03 Think of it like an onboarding document for a
new employee. It tells the agent who you are,
07:07 what the project's about, and how you want things
done. Let's build one together right now. In VS
07:14 Code, I'm creating a new file in the root of our
project. I'll name it claude.md all caps. And now,
07:21 let's fill it in section by section. Okay, the
first section is going to be project context. And
07:26 this is going to tell the agent what the workspace
is for. So I am going to give it a heading right
07:31 here at the beginning. And to do this, use the
pound key and then a space. And then you can just
07:36 type your heading. You'll notice how it turns
blue. I'm just going to give it a couple spaces
07:40 just so everything's easier to read and so I can
take a quick glance at and see what's happening.
07:46 I'm going to go and put this. So for project
context, this is my AI agent workspace. I use it
07:52 for research, content creation, and productivity
workflows. I'm going to give it a couple spaces.
07:57 Next is the about me. This is where you tell the
agent about who you are, what what you care about.
08:03 So, I'll give it again another header and give
it a couple spaces. And this is what I'm going
08:08 to put. I create content about technology and
productivity. My audience is people who want
08:12 practical, nononsense tutorials. I prefer clear,
jargon-free output. Now, the most important part,
08:19 rules. This is where you set the guardrails. These
are the instructions the agent will follow every
08:24 single time. Let's give it a heading and a couple
spaces. And my rules are going to be listed with
08:30 dashes in front of them. So you can have a clear
list just like this. So always ask clarifying
08:37 questions before starting a complex task. Show
your plans and steps before executing. Keep
08:43 reports and summaries concise, bullet points over
paragraphs. Save all output files to the output
08:48 folder and site sources when doing research. And
finally, project structure. This tells the agent

08:55 where things go. And this is where I want things.
I have my workflows. Workflow instruction files.
09:01 My output are the finished deliverables. And
resources are the reference docs and templates.
09:07 I'm just going to save this now by hitting ctrls
together. That took us about 2 minutes. But now,
09:12 every single time you open Claude Code in this
folder, it already knows your preferences, your
09:16 rules, and exactly how you want things organized.
No repeating yourself. Let me show you why this
09:22 matters so much. Now, I'm going to use the same
prompt, research remote work trends on the left
09:28 without claw.md. I made a separate folder with no
claw.md file in it. So, I get this generic wall
09:35 of text. You can see as it goes through, there's
not really much structure. I get one question just
09:40 about allowing something uh just kind of a data
dump. But on the right with the claw.md file,
09:46 the agent's asking me many different questions
as it goes through it. Then it gives me organized
09:51 bullet list, clear sections, source cited, and
save the file exactly where I told it to. Let's
09:58 take a look at this end file here. Here's the link
to the report, and it's in the output. Remember,
10:03 I told it to uh go ahead and put it the finished
output in here. So, if I click on it, it brings
10:08 me to this. Now, you might be thinking, well, that
doesn't look very good. But there's a shortcut to
10:13 see what it looks like. So, if you go to control,
shift, and v uh together, when you click in here,
10:19 if you press all those keys together, you get
the report, what it created for you. So, as I
10:24 go through it, you can see how much more detailed
uh this was. Here we have our key takeaways. we
10:31 have all of this information versus that that one
little short link that we had with the uh without
10:38 the claw.md file. And this is why the claw.md is
non-negotiable. 2 minutes of setup saves you hours
10:45 of frustration down the road. And here's something
that was just added to claude code that makes this

10:50 even more powerful. Automatic memory. As the agent works through your project, it now automatically
10:56 records things it learns and recalls them in the future sessions. So, it's just not reading your
11:01 claw.md at startup anymore. It's also building its own understanding of your project over time. The
11:07 more you use it in this folder, the smarter it gets about your preferences and how things are
11:11 organized. Before we start building our agent, I want to give you the mental model. There are
11:16 three pieces that make agendic workflows work. And once you understand them, everything else clicks
11:21 into place. Layer one, workflows. These are just plain English files that describe a process step
11:26 by step. Think of them like a recipe. The agent reads the recipe and follows it, but it's smart
11:31 enough to adapt if anything's off. Workflows live as a simple markdown file in your project. Nothing
11:37 fancy. Layer two, the agent. This is just cloud code itself. It reads your workflows,
11:42 thinks through the steps, and makes decisions. You don't program the agent. You don't write the code
11:47 for it. You just give it clear instructions and let it reason. Layer three tools. These are what
11:52 the agent uses to actually get things done. Out of the box, Claude Code can read and write files,
11:57 run terminal commands, and search the web. That's right. Web search is built right in. No extra
12:02 setup needed. The agent can go out, find current information, and bring it back into your project.
12:08 Now, you can also connect external services like Gmail, Notion, and databases using something
12:13 called MCP. That's model context protocol. That opens up a whole world of possibilities, and we'll
12:19 cover that in depth in the next video. But for today, the built-in tools are more than enough for
12:25 everything we're doing. And here's the key insight that most people miss. Most people skip straight
12:29 to the fancy tools and plugins. But the real power is writing good workflows. A well-written workflow

12:35 with basic tools will outperform a sloppy
prompt with every plug-in in the world. The
12:40 workflow is where the magic lives. Okay, enough
theory. Let's build our first agent. All right,
12:45 I'm in our empty project. You can see I just have
the claw.md file that we created. I deleted those
12:51 other test folders when I ran through that quick
test. The uh first thing we're going to do is go
12:56 and change it to plan mode. So, when I'm inside
the chat down here, I'm just going to press shift
13:01 and tab. But notice if I press it once, it goes
to edit automatically. I need to press it one
13:06 more time, it brings me to plan mode. If I press
it again, it brings me back to where I was. So,
13:11 two times pressing shift and tap brings me over
to plan mode. Now, plan mode means Cloud Code will
13:16 think and plan, but won't change any files. We're
going to design our workflow before building it.
13:21 Always plan first. Now, I'm going to describe what
I want. And watch carefully. I'm not giving it
13:27 step-by-step instructions. I'm describing my goal.
And this is what I'm going to say. I want to build
13:33 a workflow where I can give you a topic. You'll
research it thoroughly, organize the findings,
13:38 and produce a clean, structured report. Before you
start researching, you should ask me clarifying
13:42 questions about the scope, the audience, and what
I specifically want to know. I'm going to submit.
13:49 Look at this. I gave it a goal and it's already
breaking it down into steps. It's thinking about
13:55 what the workflow needs. A clarifying question
phase, a research phase, a synthesis phase,
14:01 an output phase. It's even suggesting quality
checks. That's the agendic thinking and action.
14:07 I didn't tell it these steps. It figured them out.
But I want to add one more thing before we build.
14:12 Before I add one more thing, let's go ahead and
take a look what it has already. Here we have a
14:18 resource plan workflow. And if I scroll through,
I can see there's some files that are going to be

14:23 created. And remember once we accept it to keep an
eye over on explore here because you're going to
14:27 see it get created. It's going to uh go ahead
and create this research report uh.md as I go
14:34 down. Step two, confirm the plan. Three, research.
Four, write the report. Five, save the output. It
14:39 gave me some formatting. So, if I like everything
here, I could go ahead yes and auto accept or yes,
14:45 manually approve or keep planning. But what I'm
going to say is this. I also want the report to
14:50 include a section at the end with key takeaways
and recommendations for the next step. So,
14:55 I'm just going to put that in and hit enter. So,
it's going to update the plan. It doesn't just
15:01 start over. It just folds the new requirement into
what it already had. So, that's why plan mode is
15:07 so valuable. you can shape the approach before
anything gets built. So I can see now it says
15:14 uh right here plan updated the report template
will now include two additional sections here I
15:18 have recommended next steps key takeaways. So I'm
going to go ahead and you know what I'm going to
15:24 I'm happy with it. Let's build it and I'm going
to go to auto accept. So it's just going to go
15:28 through and accept everything. So you're going
to see the mode change from plan and as I click
15:33 on this just keep an eye on explorer. So notice
we have our workflows. If I go ahead and open it,
15:40 it's still working, but we have the research
report MD that I said it was going to create.
15:46 Now let's look at what it actually wrote. This is
the important part. So I'm I'm going to go to the
15:51 research report MD file and just double click
on it. And you'll see it opens up. We have the
15:56 uh chat over here with it. But here we have the
file. Now, at the top, I've got its objective
16:02 or purpose here. A clear statement of what the
workflow does. Given a topic from the user, ask
16:08 clarifying questions. Conduct structured research.
Okay. So, let's move down to the next. We have the

16:14 qualifying question steps. Ask the user these questions. And you get an idea. What are the
16:20 specific angles of the questions you want covered? Who will read the report? Do you want a quick uh
16:25 overview or deep dive? And you can see the other ones that include uh format preferences. Now,
16:33 this is what makes it a gentic rather than just a dumb text generator. It's going to check with
16:38 you before it runs off and does the work. So, no more garbage output because it guessed wrong
16:44 about what you wanted. So, then it moves down. We can see it confirmed the plan after receiving
16:49 the answers. Confirm the research plan with users before starting. We have step three, the research
16:55 topic. as we uh go through and look at the details for each of these, you know, like avoid opinion
17:03 only posts. Uh for recent developments, include sources from the last 12 months where possible
17:08 uh what to look for while researching. Now we're on to writing the report with the writing rules,
17:15 the tone, and then saving the output. And we do have some error handling. If the topic is too
17:21 broad to research meaningfully, say so and ask the user to narrow it. Now, here's the be beautiful
17:26 part about all this. Every single word in this file is just plain English. There's no code,
17:31 no special syntax. If you wanted to change something like adding a step where it creates a
17:36 summary table or removing the clarifying questions for quick task, you just edit the text. The agent
17:43 adapts to whatever you write. So, that's the power of this approach. I want to go back over
17:48 to Claude Code and ask this question right here. What workflows do you have available? And look,
17:54 it lists our research report workflow and describes exactly what it does. It knows what
17:58 it's capable of now. So, let's put this to work. Now, for the fun part, I'm going to give it a real
18:04 research task and we're going to watch the whole agentic workflow run from start to finish. I'm

18:10 going to go and put this in. I want to research
the current state of AI agents in 2026. What are
18:15 people actually using them for? What's working?
What's overhyped? And where things are heading.
18:21 Now, watch. The first thing it should do is ask
me questions because that's what the workflow
18:26 says to do. So there it is. Look, it's asking,
you know, the audience, the scope, the depth,
18:33 all these questions. Now, I need to answer it,
but I can answer it naturally, like I'm talking
18:38 to a colleague, and this is how I'm going to do
it. Keep it broad. I want a hype versus reality
18:44 breakdown. The audience is techsavvy people
who are curious about AI, but not necessarily
18:49 developers. Medium depth. include specific tools,
platforms, stick to the last 6 to 12 months, and
18:55 formatted structured posts for clear sections and
bulleted points. So, let's submit. Now, it's going
18:60 to take those answers and start working. Watch the
to-do list. Now, here's the plan before it starts.
19:07 You can see AI agents 2026, what people are
actually doing, all the uh things I just answered,
19:13 what questions I'll research are these ones,
and I'm just going to say proceed. Yes. And
19:19 notice now even on the side we have resources a
folder that's there. I'm going to say allow. Yes.
19:30 So it's searching the web, reading through
results, pulling out relevant information.
19:35 You can see it thinking through what it finds.
And remember this is web search built right into
19:41 cloud code. No plugins, no MCP setup, no
BI API keys. It just works out of the box.
19:48 Now I have to allow it to write. All right,
it's all done. The link to the finished report
19:55 is here. It's also in my output. It lives right
here. It's the same thing. Now, the one thing I
20:00 just wanted to point out here, this right here,
the uh most honest number, even the best models
20:05 only succeed 45.7% of the time on real world task.
Human oversight isn't optional yet. I wonder how

20:12 fast this number will change. I must say it's kind
of comforting to know this. But let's go ahead and
20:17 open up the report. So remember when it opens
up the MD file, this is in markdown uh language
20:24 here. So it doesn't look quite as nice as the
formatted report. But if we hit control, shift,
20:28 and v together, it's going to give the preview of
it. So we have the topic and everything that we
20:35 said said wanted in it from the depth of 3 to five
pages, the date range, the audience. We have our
20:42 executive summary up top here, background,
uh what are the what's actually working,
20:48 key findings. So, we have this detailed report now
going all the way through uh really narrowing down
20:55 specifically what I told it to search. So,
each time I do this, it's going to be very
20:60 specific. It gives me all the sources to this.
So, the agent did all the thinking, organizing,
21:06 and all the writing. But we're not done yet. Let
me show you one more thing that makes the agendic
21:12 workflow so powerful. Iteration. I'm going to
say this. The executive summary is a bit long.
21:18 Trim it to three bullet points. Now watch. It
doesn't start over again. It goes into the report,
21:25 finds the executive summary, and trims it down.
Clean, targeted, edited. And one more thing I
21:33 wanted to do, add a comparison table of the top
five AI agent tools mentioned in the report.
21:39 So, look at that. It scans through what it already
wrote, identifies the tools it mentioned, and
21:44 created a formatted comparison table. It didn't
need me to list the tools. It pulled it from its
21:49 own research. Let's go check out those changes.
So, I'm just going to open back up this. I'll go
21:54 back to preview. We should see a shorter summary.
There's the three. 1 2 3. And now, uh, before this
22:01 was six. It didn't have quite the comparison
on it, but now 1 2 3 4 five on it. So, it made
22:08 those quick updates to the files. That's what I
mean by Agent. It remembers the context and it

22:13 knows the project and it builds on what's already there instead of starting from scratch each time.
22:18 Once you start working this way, it's really hard to go back. All right, before we wrap up, let me
22:23 save you some time with the five biggest mistakes I see people make and some pro tips that go with
22:27 them. Mistake number one, skipping claw.md. I know I keep saying it, but this one file changes
22:34 everything. Two minutes of setup saves you hours of frustration. Just do it. Mistake number two,
22:39 being too vague. Do some research gives you generic garbage, but research remote working
22:44 trends for general audience. Focus on productivity data from the last 6 months and formatted as a
22:49 structured report with bullet points. That gives you gold. Being specific is everything. Mistake
22:55 number three, not using plan mode. If you skip planning and jump right to execution, the agent
23:00 might go down the wrong path and you've wasted your time. Plan first, build second. It takes an
23:05 extra 30 seconds, but it's well worth it. Mistake number four, not telling it to ask questions.
23:11 If your workflow doesn't say ask clarifying questions before starting, the agent will assume,
23:16 and assumptions lead to wasted time and mediocre output. Always build that check-in step. Mistake
23:22 number five, trying to build everything at once. Start with one workflow, get it working well,
23:27 refine it, then build the next one. Don't try to automate your entire life in a single weekend.
23:32 Trust me on this one. Now, here are some quick pro tips. Keep your workflow files in a workflow
23:37 folder. It keeps things organized as you build more of them. Read the agents to-do list and
23:42 reasoning as it works. It's the best tool you have for understanding what's happening and catching
23:47 mistakes early before they snowball. When the output isn't right, tell the agent specifically
23:52 what to fix. Don't start over. Targeted feedback is way faster than redoing the whole thing. And

23:57 save your best claw.md and workflow files. You can copy them into new projects as starting points.
24:03 Over time, you'll build a library of workflows that make it incredibly efficient. One more,
24:08 and this is a really useful one. If the agent goes off in the wrong direction,
24:12 hover over a previous point in the thread, and you have the option to roll back the conversation,
24:17 and this undoes file changes that the agent made. So, you can rephrase your prompt and try again. No
24:23 need to manually undo edits or start over again. It's a great safety net, especially when you're
24:28 still learning how to write good prompts. And finally, effort levels. Type effort followed by
24:34 low, medium, or high to control how deeply the agent thinks before responding. Low is
24:39 fast and lightweight, great for simple edits. High means it really slows down reasons carefully and
24:44 produces its best work. Use that for more complex research or multi-step task. If you're not sure,
24:49 just leave it on the default and Claude will figure it out. But knowing this dial exists is
24:54 useful once you start running longer workflows. Let's bring it all together. In the first video,
24:58 we used Claude code to build an app. Today, we used it to build an agent. The first one does
25:04 what you tell it. The second one figures things out on its own. That's a huge jump. And that's
25:09 what you did today. So, here's my challenge for you. Build one workflow this week, just
25:13 one. Pick something you do regularly, research, writing, organizing, analysis, whatever it is,
25:19 and turn it into a workflow file. It doesn't have to be perfect. Just get it running and see what
25:24 happens. Let me know down below in the comments what you come up with. I'm always curious to see
25:28 what people are thinking about, what they could automate through one of these workflows. Thanks
25:33 for watching this time on Teachers Tech. I'll see you next week with more tech tips and tutorials.